

High Performance Computing

R. Guivarc'h

Ronan.Guivarch@toulouse-inp.fr

2025 - ModIA

Acknowledgements

This course uses materials from different sources

- ©EPCC, The University of Edinburgh, www.epcc.ed.ac.uk,
under a Creative Commons Attribution- NonCommercial-ShareAlike 4.0
International License.
http://creativecommons.org/licenses/by-nc-sa/4.0/deed.en_US



- IDRIS (Institut du développement et des ressources en informatique scientifique)
<http://www.idris.fr/formations/mpi/>
- "Calcul Réparti et Grid Computing", Patrick Amestoy and Michel Daydé (Toulouse INP - ENSEEIHT - IRIT) with J.-Y. L'Excellent (INRIA/LIP-ENS Lyon)

Introduction

Summit Overview



Compute Node

2 x POWER9
6 x NVIDIA GV100
NVMe-compatible PCIe 1600 GB SSD



25 GB/s EDR IB- (2 ports)
512 GB DRAM- (DDR4)
96 GB HBM- (3D Stacked)
Coherent Shared Memory

Components

IBM POWER9

- 22 Cores
- 4 Threads/core
- NVLink



NVIDIA GV100

- 7 TF
- 16 GB @ 0.9 TB/s
- NVLink



Compute Rack

18 Compute Servers
Warm water (70°F direct-cooled components)
RDHX for air-cooled components



39.7 TB Memory/rack
55 KW max power/rack

Compute System

10.2 PB Total Memory
256 compute racks
4,608 compute nodes
Mellanox EDR IB fabric
200 PFLOPS
~13 MW



GPFS File System

250 PB storage
2.5 TB/s read, 2.5 TB/s write



Introduction

Summit Overview



OpenMP

Components IBM POWER9

- 22 Cores
- 4 Threads/core
- NVLink



NVIDIA GV100

- 7 TF
- 16 GB @ 0.9 TB/s
- NVLink



Compute Node

2 x POWER9
6 x NVIDIA GV100
NVMe-compatible PCIe 1600 GB SSD



25 GB/s EDR IB- (2 ports)
512 GB DRAM- (DDR4)
96 GB HBM- (3D Stacked)
Coherent Shared Memory

Compute Rack

18 Compute Servers
Warm water (70°F direct-cooled components)
RDHX for air-cooled components



39.7 TB Memory/rack
55 KW max power/rack

Compute System

10.2 PB Total Memory
256 compute racks
4,608 compute nodes
Mellanox EDR IB fabric
200 PFLOPS
~13 MW



GPFS File System

250 PB storage
2.5 TB/s read, 2.5 TB/s write



Introduction

Summit Overview



Compute Node

2 x POWER9
6 x NVIDIA GV100
NVMe-compatible PCIe 1600 GB SSD



25 GB/s EDR IB- (2 ports)
512 GB DRAM- (DDR4)
96 GB HBM- (3D Stacked)
Coherent Shared Memory

Components IBM POWER9

- 22 Cores
- 4 Threads/core
- NVLink



NVIDIA GV100

- 7 TF
- 16 GB @ 0.9 TB/s
- NVLink



Compute Rack

18 Compute Servers
Warm water (70°F direct-cooled components)
RDHX for air-cooled components



39.7 TB Memory/rack
55 KW max power/rack

MPI

Compute System

10.2 PB Total Memory
256 compute racks
4,608 compute nodes
~~Memory EDR IB fabric~~
200 PFLOPS
~13 MW



GPFS File System

250 PB storage
2.5 TB/s read, 2.5 TB/s write



Introduction

Summit Overview



Compute Node

2 x POWER9
6 x NVIDIA GV100
NVMe-compatible PCIe 1600 GB SSD



25 GB/s EDR IB- (2 ports)
512 GB DRAM- (DDR4)
96 GB HBM- (3D Stacked)
Coherent Shared Memory

Components

IBM POWER9

- 22 Cores
- 4 Threads/core
- NVLink



NVIDIA GV100

- 7 TF
- 16 GB @ 0.9 TB/s
- NVLink



Compute Rack

18 Compute Servers
Warm water (70°F direct-cooled components)
RDHX for air-cooled components



39.7 TB Memory/rack
55 KW max power/rack

Compute System

10.2 PB Total Memory
256 compute racks
4,608 compute nodes
Mellanox EDR IB fabric
200 PFLOPS
~13 MW



GPFS File System

250 PB storage
2.5 TB/s read, 2.5 TB/s write



GPU programming

Introduction

Summit Overview



Compute Node

2 x POWER9
6 x NVIDIA GV100
NVMe-compatible PCIe 1600 GB SSD



25 GB/s EDR IB- (2 ports)
512 GB DRAM- (DDR4)
96 GB HBM- (3D Stacked)
Coherent Shared Memory

Components

IBM POWER9

- 22 Cores
- 4 Threads/core
- NVLink



NVIDIA GV100

- 7 TF
- 16 GB @ 0.9 TB/s
- NVLink



Compute Rack

18 Compute Servers
Warm water (70°F direct-cooled components)
RDHX for air-cooled components



39.7 TB Memory/rack
55 KW max power/rack

OpenMP

Compute System

10.2 PB Total Memory
256 compute racks
4,608 compute nodes
Mellanox EDR IB fabric
200 PFLOPS
~13 MW



GPFS File System

250 PB storage
2.5 TB/s read, 2.5 TB/s write



Part I - MPI

- ① Presentation of the SimGrid Environment
- ② Codes that illustrate some MPI concepts on Moodle

1. Message-Passing Concepts

- message passing model
- SPMD
- communication modes
- collective communications

Programming Models

Serial Programming

Concepts

Arrays	Subroutines
Control flow	Variables
Human-readable	OO

Languages

Python	C/C++
Java	Fortran
struct	if/then/else

Implementations

gcc -O3	pgcc -fast
icc	
crayftn	javac
craycc	

Message-Passing Parallel Programming

Concepts

Processes	Send/Receive
SPMD	Collectives
Groups	

Libraries

MPI
MPI_Init()


Implementations

Intel MPI	MPICH2
	Cray MPI
OpenMPI	IBM MPI

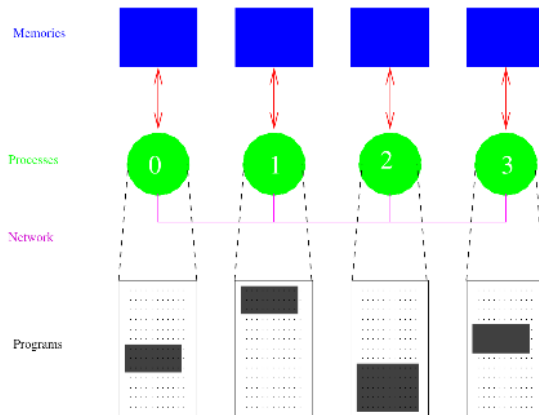
Message Passing Model

- The message passing model is based on the notion of processes - can think of a process as an instance of a running program, together with the program's data
- In the message passing model, parallelism is achieved by having many processes co-operate on the same task
- Each process has access only to its own data - ie all variables are private
- Processes communicate with each other by sending and receiving messages - typically library calls from a conventional sequential language


Sequential Paradigm




Parallel Paradigm



Distributed-Memory Architectures




Process Communication




Process Communication



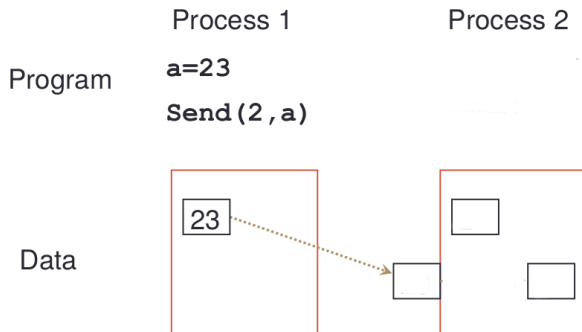
Process Communication




Process Communication




Process Communication




Process Communication




Process Communication



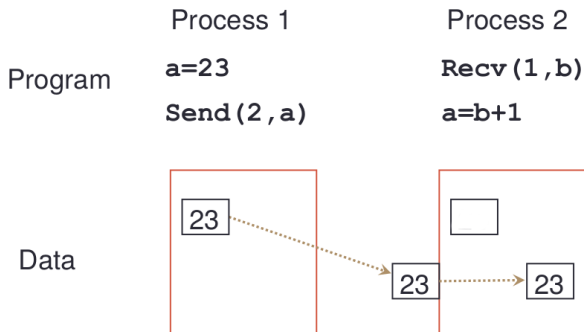
Process Communication



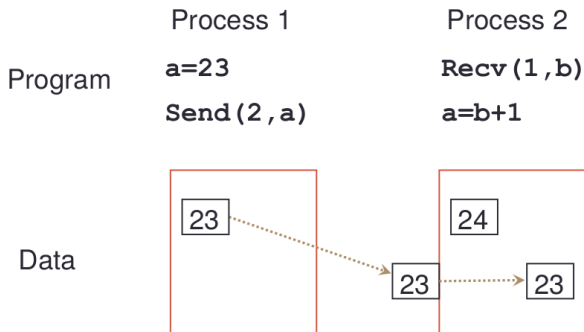
Process Communication



Process Communication




Process Communication



Flynn's taxonomy


Flynn's taxonomy is a classification of computer architectures, proposed by Michael J. Flynn in 1966 and extended in 1972.
(source: wikipedia)

- Single instruction stream, single data stream (SISD)




Flynn's taxonomy

- Single instruction stream, single data stream (SISD)
- Single instruction stream, multiple data streams (SIMD)




Flynn's taxonomy

- Single instruction stream, single data stream (SISD)
- Single instruction stream, multiple data streams (SIMD)
- Multiple instruction streams, single data stream (MISD)



Flynn's taxonomy

- Single instruction stream, single data stream (SISD)
- Single instruction stream, multiple data streams (SIMD)
- Multiple instruction streams, single data stream (MISD)
- Multiple instruction streams, multiple data streams (MIMD)



Flynn's taxonomy

Although these are not part of Flynn's work, some further divide the MIMD category into the two categories:

- SPMD, Single Program (Process), Multiple Data.
- MPMD, Multiple programs, multiple data:
"Manager/Worker" strategy, ...

SPMD

- Most message passing programs use the Single-Program-Multiple-Data (SPMD) model
- All processes run (their own copy of) the same program
- Each process has a separate copy of the data or a portion of the data
- To make this useful, each process has a unique identifier
- Processes can follow different control paths through the program, depending on their process ID
- Usually run one process per processor / core

From SPMD to MPMD

```
int main (int argc, char *argv[]) {  
  
    if (manager_process) {  
        Manager( /* Arguments */ );  
    } else {  
        Worker( /* Arguments */ );  
    }  
  
}
```


Messages


- A message transfers a number of data items of a certain type from the memory of one process to the memory of another process
- A message typically contains
 - the ID of the sending process
 - the ID of the receiving process
 - the type of the data items
 - the number of data items
 - the data itself
 - a message type identifier

Communication modes

- Sending a message can either be synchronous or asynchronous
- A synchronous send is not completed until the message has started to be received
- An asynchronous send completes as soon as the message has gone
- Receives are usually synchronous - the receiving process must wait until the message arrives


Synchronous send

- Analogy with faxing a letter (or a phone call).
- Know when fax (phone call) has started to be received.



Asynchronous send

- Analogy with posting a letter.
- Only know when letter has been posted, not when it has been received.



Point-to-Point Communications


- We have considered two processes
 - one sender
 - one receiver
- This is called point-to-point communication
 - simplest form of message passing
 - relies on matching send and receive

Collective Communications

- There are many instances where communication between groups of processes is required
- Can be built from simple messages, but often implemented separately, for efficiency


Barrier

- Global synchronisation




Broadcast

- From one process to all others




Scatter

- Information scattered to many processes




Gather

- Information gathered onto one process



Reduction

- Combine data from several processes to form a single result
- Form a global sum, product, max, min, etc.




Development of a Message-Passing Program

- Write a **single piece** of source code
 - with calls to message-passing functions such as send / receive
- Compile with a **standard compiler** and link to a **message-passing library** provided for you
 - both open-source and vendor-supplied libraries exist
- Run **multiple copies** of **same executable** on parallel machine
 - each copy is a separate process
 - each has its own private data completely distinct from others
 - each copy can be at a completely different line in the program
- Running is usually done via a launcher program
 - *"please run N copies of my executable called program.exe"*


Points of focus

- Sends and receives must match
 - danger of deadlock
 - program will stall (forever!)
- Possible to write very complicated programs, but ...
 - most scientific codes have a simple structure
 - often results in simple communications patterns
- Use collective communications where possible
 - may be implemented in efficient ways

Example of Data transfer



Example of Data transfer



2. Introduction to MPI

What is MPI?

MPI Forum

- First message-passing interface standard.
- Sixty people from forty different organisations.
- Users and vendors represented, from the US and Europe.
- Two-year process of proposals, meetings and review.
- **Message Passing Interface** document produced in 1993

History (source IDRIS)

- **Version 1.0:** June 1994, the MPI (Message Passing Interface) Forum, with the participation of about forty organisations, developed the definition of a set of subroutines concerning the MPI library.
- **Version 1.1:** June 1995, only minor changes.
- **Version 1.2:** 1997, minor changes for more consistency in the names of some subroutines.
- **Version 1.3:** September 2008, with clarifications of the MPI 1.2 version which are consistent with clarifications made by MPI-2.1.
- **Version 2.0:** Released in July 1997, important additions which were intentionally not included in MPI 1.0 (process dynamic management, one-sided communications, parallel I/O, etc.).
- **Version 2.1:** June 2008, with clarifications of the MPI 2.0 version but without any changes.
- **Version 2.2:** September 2009, with only "small" additions.

History (cont.)

- **Version 3.0:** September 2012 Changes and important additions compared to version 2.2:
 - Nonblocking collective communications
 - Revised implementation of one-sided communications Fortran (2003-2008) bindings
 - C++ bindings removed
 - Interfacing of external tools (for debugging and performance measurements)
- **Version 3.1:** June 2015
 - Correction to the Fortran (2003-2008) bindings
 - New nonblocking collective I/O routines
- **Version 4.0** was approved by the MPI Forum on June 9, 2021.
- **Version 4.1** was approved by the MPI Forum on November 2, 2023.
<https://www.mpi-forum.org>

Implementation

- MPI is a library of function/subroutine calls
- MPI is not a language
- There is no such thing as an MPI compiler
- The C/C++ or Fortran compiler you invoke knows nothing about what MPI actually does
 - only knows prototype/interface of the function/subroutine calls

Goals and Scope of MPI

- MPI's prime goals are:
 - to provide source-code portability.
 - to allow efficient implementation.
- It also offers:
 - a great deal of functionality.
 - support for heterogeneous parallel architectures.

Header files

- C/C++:

```
#include <mpi.h>
```

- Fortran 77:

```
#include 'mpif.h'
```

- Fortran 90:

```
use mpi
```

- Fortran 2008:

```
use mpi_f08
```

MPI Function Format

- C:

```
error = MPI_Xxxxx(parameter1, ...);
```

```
MPI_Xxxxx(parameter1, ...);
```

Handles

- MPI controls its own internal data structures.
- MPI defines some new types.
- MPI releases 'handles' to allow programmers to refer to these.
- C handles are of defined typedefs.

Initialising MPI

- C:

```
int MPI_Init(int *argc, char ***argv);
```

- Must be the first MPI procedure called.
 - but multiple processes are already running before MPI_Init


MPI_Init

```
int main(int argc, char *argv[]) {  
    ...  
    MPI_Init(&argc, &argv);  
    ...  
}
```

```
int main(void) {  
    ...  
    MPI_Init(NULL, NULL);  
    ...  
}
```

MPI_COMM_WORLD

- Communicators



Rank

- How do you identify different processes in a communicator?

```
int MPI_Comm_rank(MPI_Comm comm, int *rank);
```

- The rank is not the physical processor number.
 - numbering is always $\{0, 1, 2, \dots, N - 1\}$
- N is the total number of MPI process; it is given by

```
int MPI_Comm_size (MPI_Comm comm, int *size);
```

MPI_Comm_rank

```
int rank, size;  
...  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
MPI_Comm_size(MPI_COMM_WORLD, &size);  
printf("Hello from process %d of %d\n",  
       rank, size);
```

Exiting MPI

- C:

```
int MPI_Finalize(void);
```

- Must be the last MPI procedure called.

```
MPI_Finalize();
```

What machine am I on?

- Can be useful on a cluster (e.g. to confirm mapping of processes to nodes/processors/core)

```
int namelen;  
char procname[MPI_MAX_PROCESSOR_NAME];  
...  
MPI_Get_processor_name(procname, &namelen);  
printf("rank_%d_is_on_machine_%s\n",  
       rank, procname);
```

A first example

- ① (Presentation of the SimGrid Environment)
 - ② (Codes on Moodle)
 - ③ Fill the exemple "who_am_i.c" with the 5 previous MPI functions
 - ④ Compile It
 - ⑤ Run the code with different numbers of (simulated) processors
- documentation and examples: <https://rookiehpc.org>

Open source MPI implementations (IDRIS)

These can be installed on a large number of architectures but their performance results are generally inferior to the implementations of the constructors.

- **MPICH**: <http://www.mpich.org/>
- **Open MPI**: <http://www.open-mpi.org/>

MPI manpages

- documentation of distributions: `mpich`, `open-mpi`
- documentation and examples: <https://rookiehpc.org>
[example of `MPI_Comm_size`]

References (IDRIS)

- MPI : A Message-Passing Interface Standard, Version 3.1. High-Performance Computing Center Stuttgart (HLRS), University of Stuttgart, 2015. <https://fs.hlr.de/projects/par/mpi/mpi31/>
- William Gropp, Ewing Lusk, and Anthony Skjellum. Using MPI, third edition Portable Parallel Programming with the Message-Passing Interface, MIT Press, 2014.
- William Gropp, Torsten Hoefler, Rajeev Thakur and Erwing Lusk : Using Advanced MPI Modern Features of the Message-Passing Interface, MIT Press, 2014.
- Additional references:
 - Message Passing Interface Forum
<http://www.mpi-forum.org>
 - <http://www.mcs.anl.gov/research/projects/mpi/learning.html>

Tools (IDRIS)

- Debuggers
 - **Totalview**: <https://totalview.io/>
 - **Linaro Forge ex-DDT**: <https://www.linaroforge.com/>
- Performance measurement
 - **MPE** MPI Parallel Environment: <http://www.mcs.anl.gov/research/projects/perfvis/download/index.htm>
 - **FFMPI** Fast Profiling library for MPI: <http://www.mcs.anl.gov/research/projects/fpmpi/WWW/>
 - **Scalasca** Scalable Performance Analysis of Large-Scale Applications: <http://www.scalasca.org/>
- Simulation: **SimGrid**
 - <https://simgrid.org>
 - https://simgrid.github.io/SMPI_CourseWare/

Open source parallel scientific libraries

- **ScaLAPACK** Linear algebra problem solvers using direct methods:
<http://www.netlib.org/scalapack/>
- **PETSc** Linear and non-linear algebra problem solvers using iterative methods: <https://petsc.org/release/>
- **PaStiX** Parallel sparse direct Solvers (Bordeaux):
<https://solverstack.gitlabpages.inria.fr/pastix/>
- **FFTW** Fast Fourier Transform: <http://www.fftw.org>
- **FEAST** Eigenvalue Solver: <https://arxiv.org/abs/2002.04807>
- **FreeFem++** Partial Differential Equation solver <https://freefem.org/>
- **MUMPS** Parallel sparse direct Solver (Toulouse-Lyon):
<https://mumps-solver.org/>

3. Point-to-Point Communication

Point-to-Point Communication


MPI Messages

- A message contains a number of elements of some particular datatype
- MPI datatypes:
 - Basic types
 - Derived types
- Derived types can be built up from basic types.

MPI Basic Datatypes - C

MPI datatypes	C datatypes
MPI_CHAR	signed char
MPI_INT	signed int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_BYTE	char
...	...

Point-to-Point Communication



- Communication between two processes
- Source process sends message to destination process
- Communication takes place within a communicator
- Source and Destination processes are identified by their rank in the communicator

Point-to-point messaging in MPI

- Sender calls a SEND routine
 - specifying the data that is to be sent
 - this is called the ***send buffer***
- Receiver calls a RECEIVE routine
 - specifying where the incoming data should be stored
 - this is called the ***receive buffer***
- Data goes into the receive buffer
- Metadata describing message also transferred
 - this is received into separate storage
 - this is called the ***status***

Communication modes


Mode	Notes
Synchronous send	Only completes when the receive has completed
Buffered send	Always completes (unless an error occurs), irrespective of receiver \equiv Asynchronous
Standard send	Either synchronous or buffered
Receive	Completes when a message has arrived, always synchronous

MPI Communication Modes


Operation	MPI Call
Synchronous send	MPI_Ssend
Buffered send	MPI_Bsend
Standard send	MPI_Send
Receive	MPI_Recv

Synchronous Blocking Message-Passing

- Processes synchronise
- Sender process specifies the synchronous mode
- Blocking: both processes wait until the transaction has completed



Synchronous Blocking Message-Passing



- The rendezvous protocol is generally the protocol used for synchronous sends (implementation-dependent). The return receipt is optional

Sending a message with synchronous mode

- C:

```
int MPI_Ssend(void *buf,  
              int count,  
              MPI_Datatype datatype,  
              int dest,  
              int tag,  
              MPI_Comm comm);
```

Send data from rank 1 to rank 3

```
// Array of ten integers
int x[10];
...
if (rank == 1) {
    MPI_Ssend(x, 10, MPI_INT, 3, 0, MPI_COMM_WORLD);
}

// Integer scalar
int x;
...
if (rank == 1) {
    MPI_Ssend(&x, 1, MPI_INT, 3, 0, MPI_COMM_WORLD);
}
```


Receiving a message

- C:

```
int MPI_Recv(void *buf,  
             int count,  
             MPI_Datatype datatype,  
             int source,  
             int tag,  
             MPI_Comm comm,  
             MPI_Status *status);
```

Receive data from rank 1 on rank 3

```
// Array of ten integers
int y[10];
MPI_Status status;
...
if (rank == 3) {
    MPI_Recv(y, 10, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
}

// Integer scalar
int y;
MPI_Status status;
...
if (rank == 3) {
    MPI_Recv(&y, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
}
```

Implementation of synchronous send and receive

- ① synchronous send and receive of one integer (02_Ssend_Recv)
- ② synchronous send and receive of a vector of integers (03_Ssend_RecvV)

For a communication to succeed:

Synchronous or Asynchronous (see later for precision on asynchronous mode)

- Sender must specify a valid destination rank
- Receiver must specify a valid source rank
- The communicator must be the same
- Tags must match
- Message types must match
- Receiver's buffer must be large enough

4. Status, Tag, Communicators

Status, Tag, Communicators

Wildcarding

- Receiver can wildcard
- To receive from any source `MPI_ANY_SOURCE`
- To receive with any tag `MPI_ANY_TAG`
- Actual source and tag are returned in the receiver's **status** parameter

Communication Envelope Information

- Envelope information is returned from MPI_RECV as **status**
- Information includes:
 - Source: `status.MPI_SOURCE`
 - Tag: `status.MPI_TAG`
 - Count: `MPI_Get_count`

Examples with `03_Ssend_RecvV`

Received Message Count

- C:

```
int MPI_Get_count( MPI_Status *status,  
                  MPI_Datatype datatype,  
                  int *count);
```

- Example with 03_Ssend_RecvV [same size, less data, too much data]

Checking for Messages

- MPI allows you to check if any messages have arrived
 - you can "probe" for matching messages
 - same syntax as receive except no receive buffer specified

- C:

```
int MPI_Probe(int source, int tag,  
              MPI_Comm comm, MPI_Status *status);
```

- Status is set as if the receive took place
 - e.g. you can find out the size of the message and allocate space prior to receive (see RookieHPC example)
- Be careful with wildcards
 - you can use, e.g., MPI_ANY_SOURCE in call to probe
 - but must use specific source in receive to guarantee matching same message
 - MPI_Recv(buff, count, datatype, status.MPI_SOURCE, ...)

Tags


- Every message can have a tag
 - this is a non-negative integer value
 - not everyone uses them; many MPI programs set all tags to zero
- Tags can be useful in some situations
 - can choose to receive messages only of a given tag
- Most commonly used with `MPI_ANY_TAG`
 - receives the most recent message regardless of the tag
 - user then finds out the actual value by looking at the status

Communicators

- All MPI communications take place within a communicator
 - a communicator is fundamentally a group of processes
 - here is a pre-defined communicator: `MPI_COMM_WORLD` which contains ALL the processes
 - also `MPI_COMM_SELF` which contains only one process
- A message can ONLY be received within the same communicator from which it was sent
 - unlike tags, it is not possible to wildcard on `comm`

Uses of Communicators (i)

- Can split `MPI_COMM_WORLD` into pieces (example: `07_MPI_Comm_Split()`)
 - each process has a new rank within each sub-communicator
 - guarantees messages from the different pieces do not interact
 - can attempt to do this using tags but there are no guarantees



Uses of Communicators (ii)

- Can make a copy of `MPI_COMM_WORLD`
 - e.g. call the `MPI_Comm_dup` routine
 - containing all the same processes but in a new communicator
- Enables processes to communicate with each other safely within a piece of code
 - guaranteed that messages cannot be received by other code
 - this is **essential** for people writing parallel libraries (e.g. a Fast Fourier Transform, Eigenvalues computation) to stop library messages becoming mixed up with user messages
 - user cannot intercept the the library messages if the library keeps the identity of the new communicator a secret
 - not safe to simply try and reserve tag values due to wildcarding

5. Collective Communications

Collective Communications


Collective Communication

- Communications involving a group of processes
- Called by all processes in a communicator
- Examples:
 - Barrier synchronisation
 - Broadcast, scatter, gather
 - Global sum, global maximum, etc.

Characteristics of Collective Comms

- Collective action over a communicator
- All processes must communicate
- Standard collective operations are blocking
 - non-blocking versions recently introduced into MPI 3.0
- No tags
- Receive buffers must be exactly the right size


Barrier Synchronisation



C:

```
int MPI_Barrier (MPI_Comm comm);
```


Broadcast



C:

```
int MPI_Bcast (void *buffer, int count, MPI_Datatype datatype,  
               int root, MPI_Comm comm);
```


Scatter



C:

```
int MPI_Scatter(void *sendbuf, int sendcount, MPI_Datatype sendtype,
               void *recvbuf, int recvcount, MPI_Datatype recvtype,
               int root, MPI_Comm comm);
```



Gather



C:

```
int MPI_Gather(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
              void *recvbuf, int recvcount, MPI_Datatype recvtype,  
              int root, MPI_Comm comm);
```

Allgather




Global Reduction Operations

- Used to compute a result involving data distributed over a group of processes
- Examples:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation

Predefined Reduction Operations

MPI Name	Function
MPI_MAX	maximum
MPI_MIN	minimum
MPI_SUM	sum
MPI_PROD	product
MPI_MAXLOC	maximum and location
MPI_MINLOC	minimum and location
...	...

Reduce



MPI_Reduce

C:

```
int MPI_Reduce(void *sendbuf, void *recvbuf,  
               int count, MPI_Datatype datatype,  
               MPI_Op op, int root, MPI_Comm comm);
```

Example of Global Reduction

C:


```
int x, result;  
MPI_Reduce(&x, &result, 1, MPI_INT,  
           MPI_SUM, 0, MPI_COMM_WORLD);
```

- Sum of all the x values is placed in result
- The result is only placed there on process 0

Variants of MPI_REDUCE

- MPI_Allreduce no root process
- MPI_Scan "parallel prefix"

MPI_Allreduce



MPI_Allreduce

C:

```
int MPI_Allreduce(void* sendbuf, void* recvbuf,  
                  int count, MPI_Datatype datatype,  
                  MPI_Op op, MPI_Comm comm);
```

MPI_Allreduce example

C:

```
int x, result;  
MPI_Allreduce(&x, &result, 1, MPI_INT,  
              MPI_SUM, MPI_COMM_WORLD);
```

- Sum of all the x values is placed in result.
- The result is stored on every process

MPI_Scan

