

High Performance Computing

R. Guivarc'h

Ronan.Guivarch@toulouse-inp.fr

2025 - ModIA


6. Synchronous and Asynchronous Modes

Synchronous and Asynchronous Modes

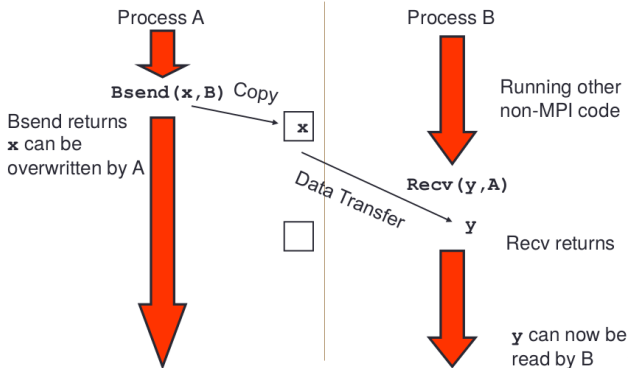
MPI Modes

- MPI_SSEND (Synchronous Send)
 - guaranteed to be synchronous
 - routine will not return until message has been delivered
- MPI_BSEND (Buffered Send)
 - guaranteed to be asynchronous
 - routine returns before the message is delivered
 - system copies data into a buffer and sends it later on
- MPI_SEND (standard Send)
 - may be implemented as synchronous or asynchronous send
 - this causes a lot of confusion (see later)

MPI_Ssend



MPI_Bsend



Bsend: notes

- Recv is always synchronous
 - if process B issued Recv before the Bsend from process A, then B would wait in the Recv until Bsend was issued
- Where does the buffer space come from?
 - for Bsend, the user provides a single large block of memory
 - make this available to MPI using `MPI_Buffer_attach`
- If A issues another Bsend before the Recv
 - system tries to store message in free space in the buffer
 - if there is not enough space then `BSEND` will FAIL!

MPI_Send

- Problems
 - Ssend runs the risk of deadlock
 - Bsend less likely to deadlock, and your code may run faster, but
 - the user must supply the buffer space
 - the routine will FAIL if this buffering is exhausted
- MPI_Send tries to solve these problems
 - buffer space is provided by the system
 - Send will normally be asynchronous (like Bsend)
 - if buffer is full, Send becomes synchronous (like Ssend)
- MPI_Send routine is unlikely to fail
 - but could cause your program to deadlock if buffering runs out

MPI_Send



- This code is **NOT** guaranteed to work
 - will deadlock if `Send` is synchronous
 - is guaranteed to deadlock if you use `Ssend`!

Examples with `06_Send_Recv`

Solutions

- To avoid deadlock
 - either match sends and receives explicitly
 - e.g. for ping-pong
 - process A sends then receives
 - process B receives then sends
- For a more general solution use non-blocking communications (see later)

7. Non Blocking Communications

Non Blocking Communications


Completion

- The **mode** of a communication determines when its constituent operations complete
 - i.e. synchronous / asynchronous
- The **form** of an operation determines when the procedure implementing that operation will return
 - i.e. when control is returned to the user program


Blocking Operations

- Relate to when the operation has completed
- Only return from the subroutine call when the operation has completed
- These are the routines you used thus far
 - MPI_Ssend
 - MPI_Bsend
 - MPI_Recv

MPI_Ssend




MPI_Bsend



Non-Blocking Operations

- Return straight away and allow the user program to continue to perform other work
- At some later time the user program can **test** or **wait** for the completion of the non-blocking operation



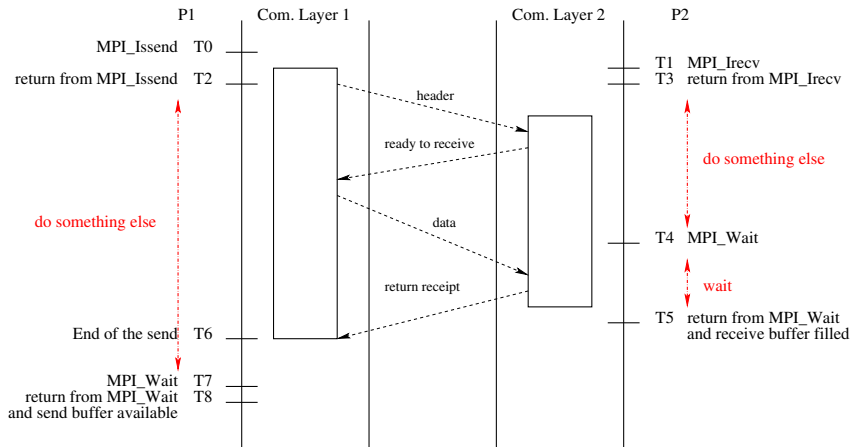
Non-Blocking Operations

- All non-blocking operations should have matching wait operations. Some systems cannot free resources until wait has been called
- A non-blocking operation immediately followed by a matching wait is equivalent to a blocking operation
- Non-blocking operations are not the same as sequential subroutine calls as the operation continues after the call has returned


Non-Blocking Communications

- Separate communication into three phases:
 - ① Initiate non-blocking communication
 - ② Do some work (perhaps involving other communications?)
 - ③ Wait for non-blocking communication to complete

Non-Blocking Communication (send and receive)



Mode vs. Form



Handles used for Non-blocking Comms

- datatype same as for blocking (`MPI_Datatype`)
- communicator same as for blocking (`MPI_Comm`)
- A request handle is allocated when a communication is initiated (`MPI_Request`)

Non-blocking Synchronous Send

"I" stands for *Immediate*

- C:

```
int MPI_Issend(void* buf, int count, MPI_Datatype datatype,
               int dest, int tag, MPI_Comm comm,
               MPI_Request *request);
```

Non-blocking Receive

- C:

```
int MPI_Irecv(void* buf, int count, MPI_Datatype datatype  
             int src, int tag, MPI_Comm comm,  
             MPI_Request *request);
```

Blocking and Non-Blocking

- Send and receive can be blocking or non-blocking
- A blocking send can be used with a non-blocking receive, and vice-versa
- Non-blocking sends can use any mode - synchronous, buffered or standard
- Synchronous mode affects completion, not initiation

Communication Modes

Non-Blocking Operation	MPI Call
Standard send	MPI_Isend
Synchronous send	MPI_Issend
Buffered send	MPI_IbSEND
Receive	MPI_Irecv

Completion

- C:

```
int MPI_Wait(MPI_Request *request,  
             MPI_Status *status);
```

```
int MPI_Test(MPI_Request *request,  
             int *flag,  
             MPI_Status *status);
```

Example (C)

```
MPI_Request request;
MPI_Status status;

if (rank == 0) {
    MPI_Issend(sendarray, 10, MPI_INT, 1, tag,
               MPI_COMM_WORLD, &request);
    Do_something_else_while_Issend_happens();
    // now wait for send to complete
    MPI_Wait(&request, &status);
} else if (rank == 1) {
    MPI_Irecv(recvarray, 10, MPI_INT, 0, tag,
               MPI_COMM_WORLD, &request);
    Do_something_else_without_data_from_recvarray
        _while_Irecv_happens();
    // now wait for receive to complete;
    MPI_Wait(&request, &status);
    Do_something_with_the_data_in_recvarray();
}
```

Multiple Communications

- Test or wait for completion of one message
- Test or wait for completion of all messages
- Test or wait for completion of as many messages as possible